

Technical Details: Engineering Practices from Side Project


Supplementary material for email to NCR Voyix CTO Author: Bojan Janjatovic, Senior Software Engineer, Platform Engineering **Date:** January 2026

Table of Contents

1. Dual-Implementation Verification
 2. Agentic AI Development with Guardrails
 3. Go for Backend Services
 4. Rust for Security-Critical Parsers
 5. Distributed Consensus for Edge Devices
 6. Emulation-Based Testing
 7. Symbolic AI for Deterministic Systems
-

1. Dual-Implementation Verification

Architecture



Test Vector Example

```
{  
  "test": "leader_election_timeout",  
  "initial_state": {  
    "role": "follower",  
    "term": 5,  
    "votes_received": 0,  
    "election_timeout_ms": 150  
  },  
  "input": {  
    "event": "timeout_elapsed",  
    "elapsed_ms": 200  
  },  
  "expected_output": {  
    "role": "candidate",  
    "term": 6,  
    "voted_for": "self",  
    "action": "broadcast_vote_request"  
  }  
}
```

Why This Works

- **Different failure modes:** C bugs (buffer overflows, pointer errors) differ from Rust bugs (logic errors, type mismatches)
 - **AI cross-check:** If AI generates C code with a subtle bug, it's unlikely to generate the identical bug in Rust
 - **Specification as code:** JSON vectors serve as executable documentation
-

2. Agentic AI Development - Two Modes

Not all tasks are equal. I distinguish between two fundamentally different modes of AI-assisted development:

Mode A: Extreme Agentic (Autonomous)

EXTREME AGENTIC MODE

When to use:

- Well-defined implementation task
- Clear success criteria exist (test vectors, specs)
- Correctness is automatically verifiable

How it works:

1. Human defines task and success criteria
2. AI works autonomously for 1+ hours
3. No human input during session
4. AI commits to git, runs tests, iterates

5. Human reviews RESULT, not every step

Verification:

- Test vectors pass in both C and Rust → implementation correct
- Emulation tests pass → behavior correct
- Human reviews final diff, not intermediate steps

Example tasks:

- "Implement heartbeat protocol per this spec"
 - "Add thermal derating per this formula"
 - "Parse CAN message ID 0x300 per this format"
-

Mode B: Human-in-the-Loop (Guided)

HUMAN-IN-THE-LOOP MODE

When to use:

- Exploratory work, research
- Architecture decisions
- No clear test criteria
- Correctness cannot be automatically verified

How it works:

1. Human and AI collaborate step-by-step
2. Human reviews each proposed change
3. Explicit permission grants per action
4. Slower, but necessary for ambiguous tasks

Example tasks:

- "Help me design the state machine"
 - "What's the best approach for X?"
 - "Refactor this module for clarity"
 - "Write documentation for this feature"
-

Decision Matrix

Can correctness be automatically verified? (test vectors, specs)	YES → EXTREME AGENTIC (autonomous, fast)
	NO → HUMAN-IN-THE-LOOP (guided, careful)

Key Insight

Test vectors are the enabler of autonomous AI development. Without them, human oversight is mandatory. The dual-implementation approach (C + Rust + shared vectors) creates the verification infrastructure that makes extreme agentic mode safe.

TEST VECTORS → AUTONOMOUS AI POSSIBLE
NO TEST VECTORS → HUMAN MUST GUIDE

This matches how we think about other automation: automated testing enables continuous deployment; without tests, manual QA is required.

3. Go for Backend Services

No Magic Programming

```
// Go: Explicit error handling - no hidden control flow
func (s *Simulation) AddModule(id uint8) error {
    if id >= MaxModules {
        return fmt.Errorf("module ID %d exceeds max %d", id, MaxModules)
    }
    if _, exists := s.modules[id]; exists {
        return ErrModuleAlreadyExists
    }
    // ...
    return nil
}
```

What “no magic” means in practice: - No exceptions flying up the stack invisibly - No inheritance hierarchies to trace - No decorators, annotations, or runtime reflection surprises - No dependency injection frameworks with hidden wiring - Code reads top-to-bottom, left-to-right

When debugging at 3 AM, “boring” code is a feature.

Docker / Container Benefits

```
# Go Dockerfile - entire thing
FROM scratch
COPY myservice /myservice
ENTRYPOINT ["/myservice"]
```

Container Image Size Comparison

Technology	Image Size	Notes
Java (Spring Boot)	400-800 MB	JVM + dependencies
Python (Flask)	150-400 MB	Interpreter + packages
Node.js (Express)	150-300 MB	Runtime + node_modules

Technology	Image Size	Notes
Go	10-20 MB	Static binary, no runtime

Why this matters: - Faster deployments (10 MB pulls faster than 400 MB) - Smaller attack surface (no runtime = fewer CVEs) - FROM scratch possible (empty base image) - No runtime dependencies - Millisecond startup (no JVM warmup)

Go Benefits Summary

Aspect	Go Advantage
Error handling	Explicit, in-band, no hidden flow
Inheritance	None - composition only
Dependencies	Single binary, zero runtime deps
Containers	10 MB images, FROM scratch
Startup	Milliseconds
Concurrency	Built-in race detector
Learning curve	Small language, readable in days

4. Rust for Security-Critical Parsers

Code Example

```
// Rust: Compiler enforces handling all cases
enum ModuleState {
    Offline,
    Initializing { start_time: Instant },
    Online { neighbor_count: u8 },
    Fault { code: FaultCode, recoverable: bool },
}

match module.state {
    ModuleState::Offline => self.start_discovery(),
    ModuleState::Fault { recoverable: true, .. } => self.attempt_recovery(),
    ModuleState::Fault { recoverable: false, .. } => self.halt(),
    // Compiler error if cases not exhaustive
}
```


Rust Benefits

- No null pointer dereferences (Option instead)
 - No data races (borrow checker enforces at compile time)
 - Pattern matching forces handling all cases
 - Memory safety without garbage collection
-


5. Distributed Consensus for Edge Devices

Architecture Comparison

TRADITIONAL



DISTRIBUTED



Single point of failure.


Any device can lead.
Others continue if one fails.

Key Insight

The same Raft consensus used in etcd/Consul works for coordinating physical devices, not just database replicas.

6. Emulation-Based Testing

Renode Architecture



Benefits

- Test firmware without hardware
 - Reproducible CI/CD pipeline
 - Catch “works on my machine” issues
 - Test failure scenarios safely
-

7. Symbolic AI for Deterministic Systems

LLM vs Safety-Critical Requirements

LLM CHARACTERISTICS

X Non-deterministic	Same input can produce different output
X No guarantees	"Usually correct" ≠ "always correct"
X Black box	Cannot prove why a decision was made
X Hallucinations	Confidently wrong outputs
X No formal verification possible	

REQUIRED FOR SAFETY-CRITICAL SYSTEMS

✓ Deterministic	Same input → same output, always
✓ Provable	Can formally verify correctness
✓ Explainable	Clear audit trail of decisions
✓ Bounded behavior	Known limits, predictable edge cases
✓ Certifiable	Meets ISO 26262, IEC 61508, etc.

Symbolic AI Approach

RULE ENGINE EXAMPLE

```
IF voltage > 920V AND current > 0
    THEN fault(OVERVOLTAGE), action(DISCONNECT)
```

```
IF temperature > 85°C AND power > 2.5kW
    THEN action(DERATE), set_power(2.0kW)
```

```
IF heartbeat_timeout > 500ms
    THEN state(SUSPECT), action(PROBE)
```

Results:

- DETERMINISTIC - Same conditions → same actions
- EXPLAINABLE - "Why did it disconnect?" → logged rule
- TESTABLE - Every rule has explicit test cases
- CERTIFIABLE - Formal methods can prove properties

Automotive Certification Example (ISO 26262)

Requirement: Demonstrate that software behaves correctly under ALL conditions

- **LLM approach:** "It worked in 99.7% of test cases" → NOT CERTIFIABLE
- **Symbolic approach:** "Here is formal proof that condition X always produces action Y" → CERTIFIABLE

Hybrid Approach: LLM + Symbolic

LLM (Claude, GPT)	SYMBOLIC LAYER
<ul style="list-style-type: none">• Development assistance• Code generation• Documentation• Test case ideation	<ul style="list-style-type: none">• Runtime decisions• Safety interlocks• Fault handling• State machine logic
Used at DESIGN TIME (human in the loop)	Used at RUN TIME (deterministic execution)

My project uses both: - LLM (Claude) helps write code, explore designs, generate documentation - Symbolic logic (state machines, explicit rules) runs on the actual device - The LLM never runs in production - only deterministic code does

Expert Systems Renaissance

Era	Approach	Limitation
1980s	Expert systems	Knowledge acquisition bottleneck
2020s	LLMs	Non-deterministic, can't certify
Now	Hybrid	LLM helps BUILD expert systems, symbolic logic RUNS them

Contact

Bojan Janjatovic Senior Software Engineer, Platform Engineering NCR Voyix

Available for deeper discussion on any of these topics.